

ALMA MATER STUDIORUM UNIVERSITÀ DI BOLOGNA  
DEPARTMENT OF ELECTRICAL, ELECTRONIC AND INFORMATION ENGINEERING  
MASTER'S DEGREE IN AUTOMATION ENGINEERING

# **A Deep Reinforcement Learning approach based on policy gradient for Mobile Robot Navigation**

*Candidate:*

Enrico Pianazzi

*Advisor:*

Prof. Giuseppe Notarstefano

*Co-Advisors:*


Dott. Andrea Camisa

Ing. Lorenzo Sforzi

# Motivations

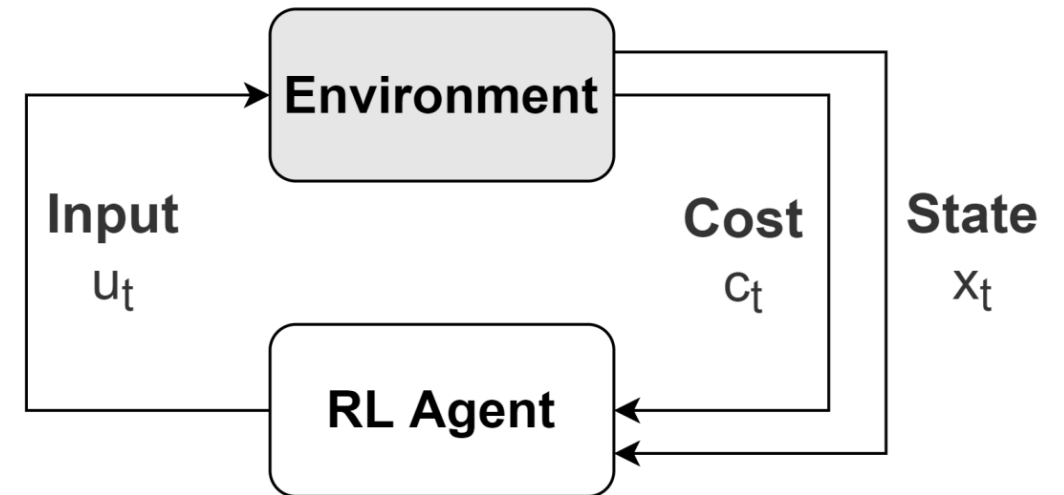
---

- Autonomous navigation features in mobile robots are highly valuable
- Reinforcement Learning can achieve complex data-driven control
- State feedback control without the need of complex sensory data



# Reinforcement Learning

- Optimize a policy (controller)  
 $\mu(x_t) = u_t$  to minimize future cost
- Only experience cost and state signals
- Division in discrete time steps and episodes



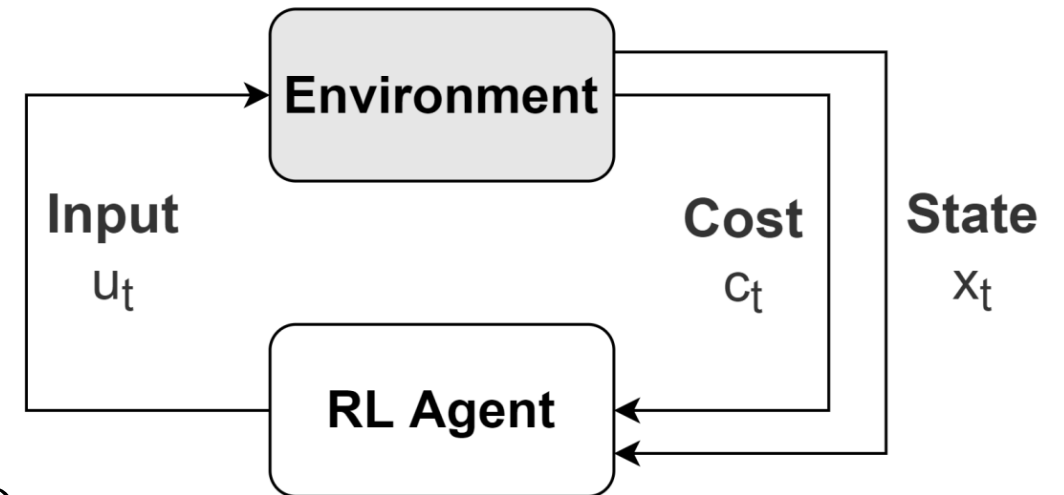
# Q-function and Bellman's equation

- Q-function: estimates future cost
- Bellman's equation:

$$Q(x_t, u_t) = c(x_t, u_t) + \gamma Q(x_{t+1}, \mu(x_{t+1}))$$

- Temporal Difference error:

$$TD_{er} = c(x_t, u_t) + \gamma Q(x_{t+1}, \mu(x_{t+1})) - Q(x_t, u_t)$$





# Problem Statement

- Task: reach goal area from any initial configuration in the environment
- Environment limits, obstacles are unknown
- Continuous states and inputs
- Robot simulated as a unicycle, but model is unknown


State:  $x = [p_x \ p_y \ \theta]$

$$\text{Model: } \dot{x} = \begin{bmatrix} \cos(\theta) & 0 \\ \sin(\theta) & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} v \\ \omega \end{bmatrix}$$



# Deep Deterministic Policy Gradient (DDPG)

- Policy: controller
- Q-function: measures the cost of a state-input pair
- Implemented as deep neural networks
- Experience Replay: keeps a buffer of past transitions {state, input, cost, next state}



## Replay Buffer


$\{x_1, u_1, c(x_1, u_1), x_2\}$

$\{x_2, u_2, c(x_2, u_2), x_3\}$

...

$\{x_\tau, u_\tau, c(x_\tau, u_\tau), x_{\tau+1}\}$

# Policy Training





- Trained to choose inputs that minimize cost
- Policy network parameters:  $\theta_k^\mu$
- Policy gradient update:

$$\theta_{k+1}^\mu = \theta_k^\mu - \alpha_\mu \nabla_{\theta^\mu} Q(x_t, \mu(x_t))$$

# Q-function Training

- Sampled transition:  $\{x_t, u_t, c(x_t, u_t), x_{t+1}\}$
- Q-function network parameters:  $\theta_k^Q$
- Gradient descent update minimizing the Temporal Difference error:

$$\theta_{k+1}^Q = \theta_k^Q - \alpha_Q \nabla_{\theta^Q} \underbrace{[c(x_t, u_t) + \gamma Q(x_{t+1}, \mu(x_{t+1}))]}_{\text{Target of the update}} - \underbrace{Q(x_t, u_t)}_{\text{Prediction}}]^2$$





# Target Networks

- Copies of the main networks:  $Q'(x, u)$  and  $\mu'(x)$ , with parameters:  $\theta^{Q'}$  and  $\theta^{\mu'}$
- Slowly track the main networks parameters:

$$\begin{cases} \theta^{Q'} = \tau \theta^Q + (1 - \tau) \theta^{Q'} \\ \theta^{\mu'} = \tau \theta^{\mu} + (1 - \tau) \theta^{\mu'} \end{cases}$$

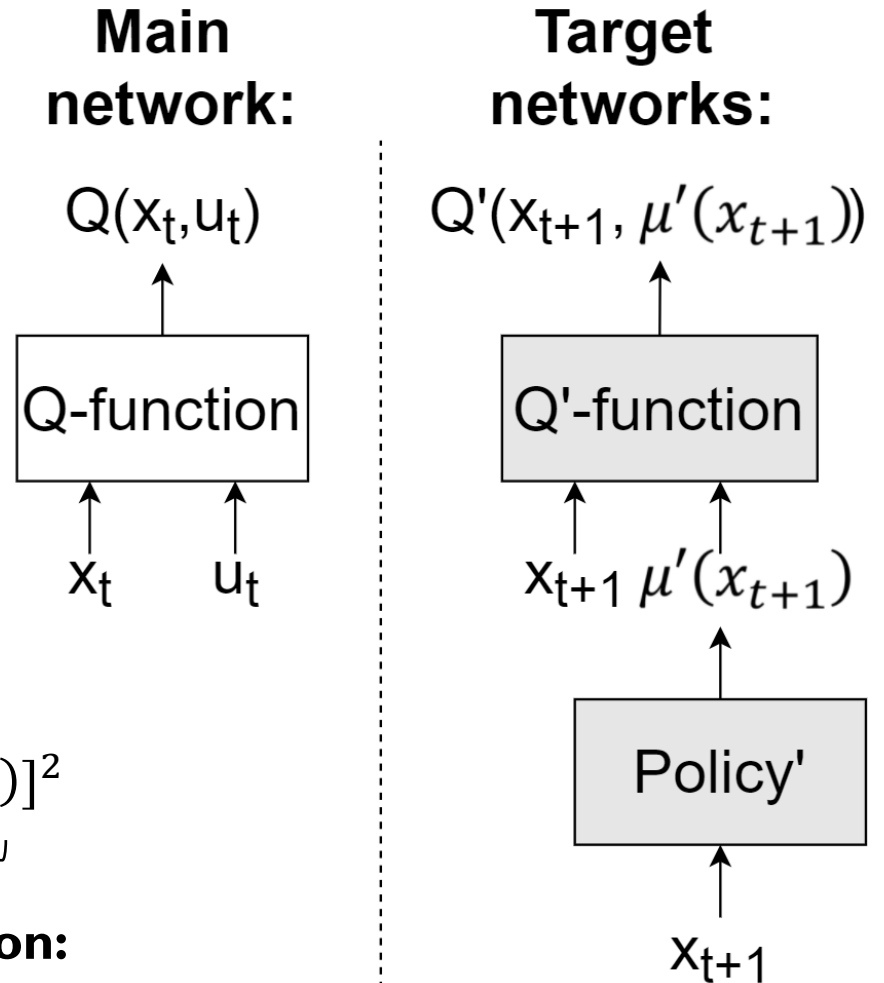
With  $\tau \ll 1$

- Only used for target computation of the TD error:

$$\theta_{k+1}^Q = \theta_k^Q - \alpha_Q \nabla_{\theta^Q} [ \underbrace{c(x_t, u_t) + \gamma Q'(x_{t+1}, \mu'(x_{t+1}))}_{\text{Target of the update: Target networks}} - \underbrace{Q(x_t, u_t)}_{\text{Prediction: main network}} ]^2$$

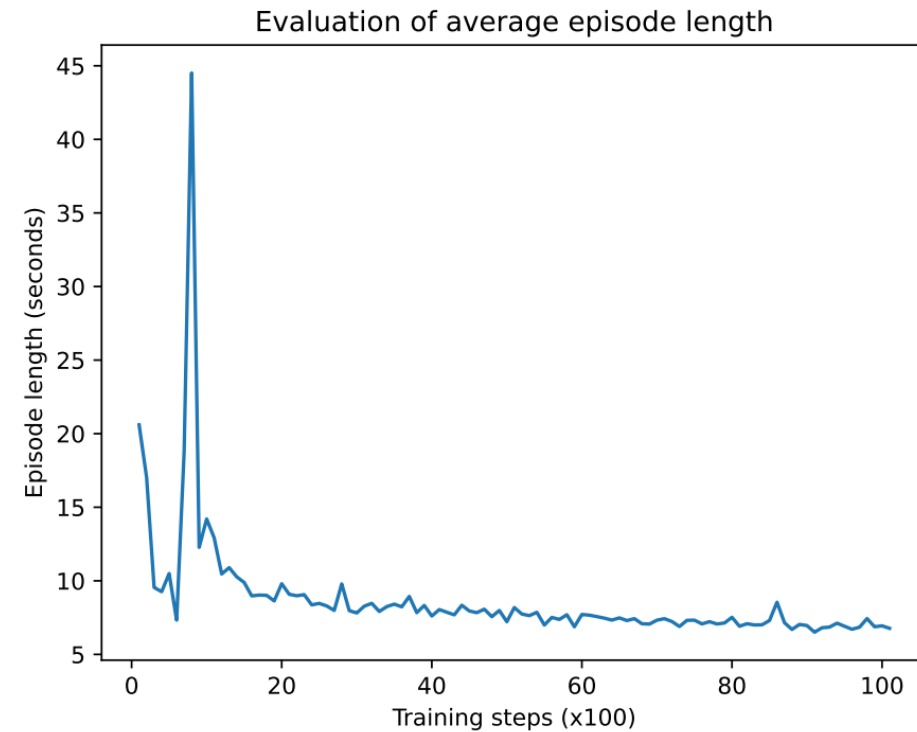
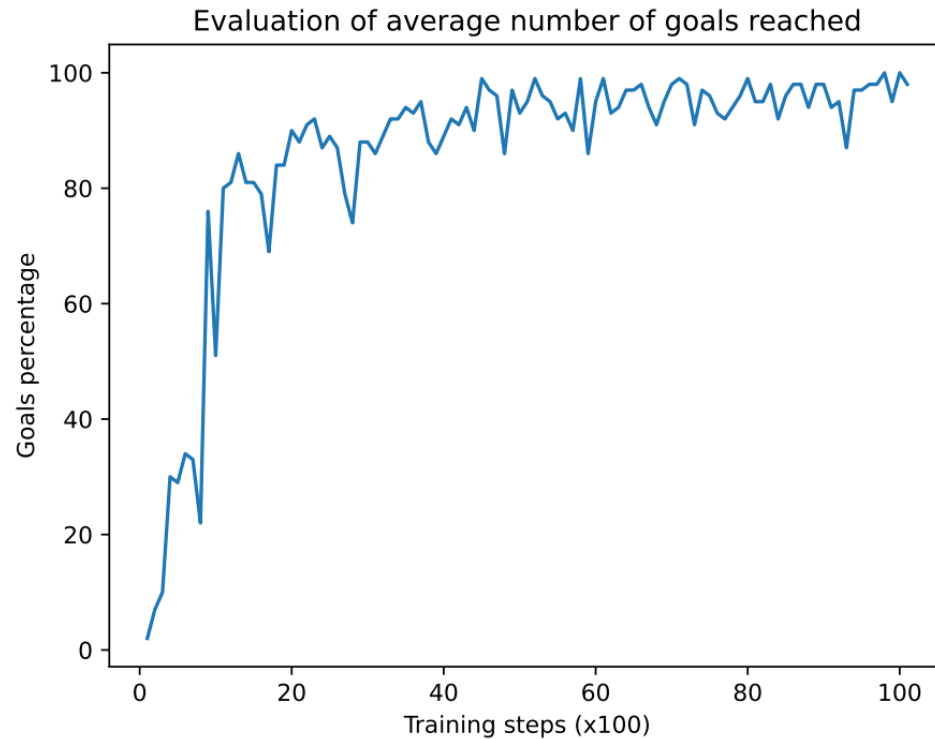
**Target of the update:**  
**Target networks**

**Prediction:**  
**main network**



# Controller Performance Evolution


- Experiments from a set of fixed initial conditions
- 10000 episodes training




# Resulting Robot Trajectories

---

10Hz simulation:



100Hz simulation:



# Conclusions

- Achieved above 95% success rate in the autonomous goal-reaching task
- Provided a custom Python implementation of DDPG and the unicycle simulation
- Drawback: length of training
- Future work: algorithm tuning; distributed setting; integration with model-based techniques